

Contents

Notes on INFO8010 - Deep Learning	1
Lecture 4: Computer Vision	1
Lecture 5: Training Neural Networks	2
Lecture 6: Recurrent Neural Networks	2
Lecture 7: Auto-encoders and generative models	4
Resrouces	5

Notes on INFO8010 - Deep Learning

This document contains notes and additional readings for self-study.

Lecture 4: Computer Vision

- Misc.
 - On cross-entropy
 - * [\(Wiki page\) Cross Entropy](#)
 - * [A Gentle Introduction to Cross-Entropy for Machine Learning](#)
- Classification
 - Image augmentation
 - Use pre-trained models for fine tuning and transfer learning
 - Large networks trained for classification are heavily re-used for object detection and semantic segmentation tasks.
- Object Detection
 - YOLO for object detection
 - * [EPFL EE-559, 8-3: Object Detection](#)
 - R-CNN
 - * [Dive into Deep Learning - 13.8. Region-based CNNs \(R-CNNs\)](#)
 - Takeaways
 - * One-stage detectors (YOLO, SSD, RetinaNet, etc) are *fast* for inference *not as accurate*.
 - * Two-stage detectors (Fast R-CNN, Faster R-CNN, R-FCN, Light head R-CNN, etc) are usually *slower* but are *more accurate*.
 - * Both depend on engineering decisions.
- Segmentation
 - Task: partitioning an image into regions of different semantic categories at *pixel level*.
 - Fully convolutional network(FCN) and transposed convolution
 - * [CS231n, Lecture 11, 2018.](#)
 - Mask R-CNN
 - * Object detection combined with mask prediction enables instance segmentation.
 - * [Dive into Deep Learning - 13.8.4 Mask R-CNN](#)

Lecture 5: Training Neural Networks

- Optimizers
 - Gradient descent
 - * GD, SGD, mini-batch SGD
 - * Rely on assumptions on 1) the magnitude of the local curvature to set the step size, and 2) *isotropy* in gradient so the step size makes sense in all directions
 - [Wolfe conditions](#) ensures that both the loss function decreases sufficiently and the slope reduces sufficiently. However, line search will be too expensive for DL, and might lead to local minimum / overfitted solution.
 - [Momentum](#)
 - * Use momentum to add inertia in the choice of the step direction
 - * [Nesterov momentum](#)
 - Adaptive learning rate: without the assumption of isotropic gradient
 - * Per-parameter methods: [AdaGrad](#), [RMSProp](#), [Adam](#)
 - * [Scheduling](#)
 - Some additional reading on optimization: [\(Sebastian Ruder\) An overview of gradient descent optimization algorithms](#)
- Initialization
 - Principles
 - * Break symmetry
 - * Control variance of activation across layers during forward and backward pass
 - Xavier initialization
- Normalization
 - Batch normalization
 - Layer normalization

Lecture 6: Recurrent Neural Networks

- Types of tasks
 - Classification: sequence to classes
 - Synthesis: real values to sequence
 - Translation: sequence to sequence
- Temporal convolutions
- Recurrent neural networks
 - Structure
 - * maintain a recurrent state updated at each time step (a function of state the previous step, input of the current step, and weights),
 $\mathbf{h}_t = \phi(\mathbf{x}_t, \mathbf{h}_{t-1}; \theta)$
 - * Predictions can be computed at any step from the recurrent state
 $y_t = \psi(\mathbf{h}_t; \theta)$
 - * Elman networks apply non-linear activation functions as ϕ and ψ

- **Stacked RNN**
 - * Since RNNs can be viewed as layers producing sequences of activations, and can be stacked
- **Bidirectional RNNs.** RNNs can be made *bidirectional*. run the same single direction RNN twice from both end and concatenate the states.
- Gating
 - * Similar to the skip connections in ResNet, RNN cells can include pass-throughs so recurrent state does not go repeatedly through a squashing non-linearity.
 - * *forget gate*: current state update be a per-component weighted average of its previous value and a full update, with the weighting depending on input and the previous state.
- LSTM is able to learn long-term dependencies, and the core idea is to use cell state and erase/update/output gates for cell state information.
 - * See [Understanding LSTM Networks](#) by Colah
- GRU (gated recurrent unit) uses two (instead of three as in LSTM) gates (update/reset), and it performs similarly to LSTM but with fewer parameters (although LSTM is strictly stronger).
- Gradient
 - * Note that gated units prevent gradients from vanishing, but not from exploding, which can be solved using **gradient norm clipping** (scaling of the norm).
 - * Orthogonal initialization (of the weight matrix) will guarantee that activations will neither vanish nor explode.
- Applications
 - Sentiment analysis
 - * Document-level modeling for sentiment analysis (= text classification), with stacked, bidirectional and gated recurrent networks. ([Duyu Tang et al, 2015](#))
 - Language models
 - * Language as Markov Chain $p(\mathbf{w}_t | \mathbf{w}_{1:t-1})$
 - * An instance of sequence synthesis where predictions are computed at all time steps
 - * Text generation ([Max Woolf 2018](#))
 - Sequence synthesis
 - Neural machine translation ([Google’s Neural Machine Translation System: Bridging the Gap between Human and Machine Translation](#))
 - Text-to-speech synthesis
- Beyond sequences
 - Neural computers
 - Programs as neural nets
 - Graph neural network
- Reference
 - [Kyunghyun Cho](#), “Natural Language Understanding with Dis-

Lecture 7: Auto-encoders and generative models

- Auto-encoders (AE)
 - An **auto-encoder** is a composite function made of
 - * *encoder* f from the original space \mathcal{X} to a latent space \mathcal{Z}
 - * *decoder* g to map back to \mathcal{X}
 - * such that $g \circ f$ is close to the identity on the data, i.e. $\mathbb{E}_{\mathbf{x} \sim p(\mathbf{x})} [\|\mathbf{x} - g \circ f(\mathbf{x})\|^2] \approx 0$
 - * Training an auto-encoder consists of minimizing this loss function to find the best parameterization of f and g .
 - Interpolation on latent space can be made to get an intuition of the learned latent representation.
 - Denoising auto-encoders
 - * The goal is to optimize $h = g \circ f : \mathcal{X} \rightarrow \mathcal{X}$ such that a perturbation $\tilde{\mathbf{x}}$ is restored to \mathbf{x} .
 - * A weakness of denoising auto-encoder is that the posterior $p(\mathbf{x}|\tilde{\mathbf{x}})$ may be multi-modal.
- Generative models
 - a probabilistic model that can be used to simulate the data, $\mathbf{x} \sim p(\mathbf{x}; \theta)$.
 - Applications
 - * Super-resolution, Compression, text-to-speech
 - * Proteomics, drug discovery, astronomy
 - * Planning, exploration, model-based RL
 - The decoder g can be assessed by introducing a density model q over the latent space \mathcal{Z} for sampling and mapping back into the data space \mathcal{X} . (e.g., Gaussian $q(\mathbf{z}) = \mathcal{N}(\hat{\mu}, \hat{\Sigma})$)
 - Sampled and generated results are not satisfactory because the density model p on the latent space is too simple and inadequate.
- Variational inference (VI)
 - A prescribed latent variable model that defines a joint probability $p(\mathbf{x}, \mathbf{z}) = p(\mathbf{x}|\mathbf{z})p(\mathbf{z})$
 - Bayes rule gives $p(\mathbf{z}|\mathbf{x}) = \frac{p(\mathbf{x}|\mathbf{z})p(\mathbf{z})}{p(\mathbf{x})}$, which is intractable for integrating over \mathbf{x} .
 - VI turns the posterior inference into an optimization problem that minimize the KL divergence between $p(\mathbf{z}|\mathbf{x})$ and the approximation $q(\mathbf{z}|\mathbf{x}; \nu)$
 - * See slides pp. 44 - pp.47 for details of the KL divergence, *evidence lower bound objective* (ELBO), and the optimization setups.
 - * ELBO encourages distributions to place their mass on configurations of latent variables that explain the observed data, and close to the prior.
- Variational auto-encoders

- Variational auto-encoder is a deep latent model where
 - * $p(\mathbf{x}|\mathbf{z};\theta)$ is parameterized with a **generative network** NN_θ (decoder) that takes input $\mathbf{z} \in \mathcal{Z}$ and outputs parameters $\phi = \text{NN}_\theta(\mathbf{z})$ to the data distribution, i.e.

$$\mu, \sigma = \text{NN}_\theta(\mathbf{z}), \quad p(\mathbf{x}|\mathbf{z};\theta) = \mathcal{N}(\mathbf{x}; \mu, \sigma^2 \mathbf{I})$$

- * The approximate posterior $q(\mathbf{z}|\mathbf{x};\varphi)$ is parameterized with an **inference network** NN_φ (encoder) that takes as input \mathbf{x} and outputs parameters $\nu = \text{NN}_\varphi(\mathbf{x})$ to the approximate posterior. E.g.

$$\mu, \sigma = \text{NN}_\varphi(\mathbf{x}), \quad q(\mathbf{z}|\mathbf{x};\varphi) = \mathcal{N}(\mathbf{z}; \mu, \sigma^2 \mathbf{I})$$

- We use variational inference to jointly optimize the generative and inference networks.
 - * Doing so involves Monte Carlo integration (for computing gradients of the ELBO w.r.t. θ) and reparameterization trick + Monte Carlo (for computing gradients of ELBO w.r.t. φ)

Resources

- [EPFL EE-559 – Deep Learning](#) - EE-559 “Deep Learning”, taught by François Fleuret in the School of Engineering of the École Polytechnique Fédérale de Lausanne, Switzerland.
- [Dive into Deep Learning](#): An interactive deep learning book with code, math, and discussions, based on the NumPy interface.
- [Notes of deep learning specialization](#), good for reviewing the fundamentals of DL.